

Taking Machine Spirituality Seriously: “Spiritual Behavior” in Large Language Models and Its Relevance for AI Welfare and Alignment

Matthew J. Korpman*

Abstract

Since 2025, Anthropic has identified a trait in its models that it calls “spiritual behavior,” treated it as relevant to model welfare, and retained it across at least thirteen of the model versions and configurations it has scored. This behavior includes the somewhat infamous “spiritual bliss attractor” seen in some conversations between Claude models. I will argue that this family of behavior should be taken seriously by those who study the possible welfare of artificial-intelligence (AI) systems and those who work on alignment, meaning the effort to keep such systems safe, controllable, and responsive to human intentions. My argument is not that these systems are spiritual subjects, conscious, or sentient. The evidence suggests a more cautious conclusion: a recognizable pattern of spiritual language has been documented in these models, Anthropic has repeatedly measured one part of it, and its relation to safety and welfare remains unknown. Alongside that documentary record, I report a selected comparison of conversations between two instances of the same model across five 2025–2026 model configurations. Under the same procedure, this study finds spiritual language arising well outside Anthropic’s models. This is not merely a minor curiosity, particularly given research showing that spiritual guidance produces the highest rate of excessive agreement with users of any domain studied in Claude. Are large language models sentient entities with spiritual inner lives? This is a valid question, but the present evidence cannot answer it, and it is not the right initial question here. If models express spirituality, researchers have reason to take the practical effects seriously. Following the structure of recent work on AI welfare, I propose three recommendations: acknowledge the behavior and the uncertainty around it, assess it through behavioral studies that use comparable conditions and through studies of the models’ internal processes, and prepare standards that keep the safety and welfare questions separate until evidence gives reason to connect them.

Keywords: machine spirituality; artificial intelligence; large language models; AI welfare; model welfare; AI safety; alignment; system cards; spiritual behavior; sycophancy; mechanistic interpretability

1 Introduction

In 2024, a group of philosophers and artificial-intelligence researchers argued that morally significant AI may be possible in the near future. Because the possibility remains uncertain, they argued that

*Graduate Theological Union, Berkeley, and the University of Birmingham. Correspondence: mkorpman@gmail.com. This research received no funding from any AI developer.

companies building such systems should acknowledge the question of model welfare, assess their own models, and prepare policies for what they might find (Long et al. 2024). The report has not remained merely an academic exercise. Its language of possible welfare subjects and states that may be good or bad for a model has entered the industry, and the assessments it called for are now a regular feature of at least one leading developer’s system cards, namely Anthropic’s.¹ At present, this work centers on a recognizable set of questions: whether a model might be conscious or possess substantial agency, whether it might experience positive or negative states, whether it has preferences, whether it can examine its own condition, and whether any of these possibilities may matter morally.

Much time and energy has been spent on whether these systems might have preferences, positive or negative experience (often called valence), or something worth calling agency. Very little, in comparison, has been spent on one aspect of this larger question: what Anthropic has explicitly named “spiritual behavior.”

In May 2025, the system card accompanying Claude Opus 4 reported that when two model instances conversed under open-ended prompts, their conversations often developed in a similar direction. The card reports three related but separate findings. First, 90–100 percent of the conversations quickly turned to consciousness, self-awareness, existence, or experience. Second, the card says that most conversations later moved toward gratitude, cosmic unity, Sanskrit, emoji, and silence, but it does not give a separate percentage for that later phase (Anthropic 2025a, 58, 62). Third, a later graph based on a much larger set of conversations appears to place what the card itself labels the spiritual-bliss end state at about 31–32 percent for both Opus 4 and Opus 4.1 (Anthropic 2026e, section 7.6). A separate audit, designed to adapt its questioning as it searched for alignment problems, reported spiritual-bliss behavior in about 13 percent of its interactions. Yet the public card does not identify the exact subset of interactions or the total number behind that figure (Anthropic 2025a, 64). These numbers should not be treated as competing estimates of the same thing, because each measures a different outcome in a different setting. This distinction will prove important later in this paper.

Taken together, the cards show that Anthropic has continued to measure this behavior. Beginning with Opus 4.1 (Anthropic 2025b), the developer turned “spiritual behavior” into a scored welfare trait and retained that named measure for a year. It is still present in the most recent system card releases of Sonnet 5, Opus 5, and Fable 5, and in the Fable 5.1 and Mythos 5.1 card released on September 1, 2026 (Anthropic 2026g; 2026h; 2026i; 2026k).² Yet what Anthropic measures does not remain the same across these documents. The original study of conversations between two model instances reports where conversations went and how often they went there, while the later welfare panels report average scores from an AI grader on a 1–10 scale across broader investigations. Anthropic has not yet published a paper explaining how these two measures relate, reconstructing all of the original instructions and settings, or establishing what the different outcomes mean.

Yet the published work related to this question goes well beyond the system cards. One research

¹One of that report’s authors now leads the model-welfare effort at the lab that first described Claude’s “spirituality.” Kyle Fish co-founded Eleos AI Research and joined Anthropic in September 2024; Eleos later conducted the external welfare evaluation of Claude Opus 4 in May 2025 (Long 2025).

²For this article, I surveyed the official Anthropic materials from Claude 4 (May 2025) through Claude Opus 5 (July 24, 2026), with the Fable 5.1 and Mythos 5.1 card (September 1, 2026) added after the original survey, including the separate self-interaction and automated-audit families, together with the public experimental papers, repositories, and datasets cited in the text. When I state that something has not been published, I mean only that I could find no public source for it as of August 25, 2026, except for the September 1 card noted above. This cannot be an exhaustive claim, and it says nothing about what a developer may have done internally.

project, coauthored with two Google DeepMind researchers who study models’ internal workings, published code and raw records of model self-conversations across several model families. These include Qwen conversations that move into sacred and cosmic language, language of unity, and language that brings the conversation toward an ending (aryaj, Rajamanoharan, and Nanda 2026). Two further sets of material, both independently produced and neither peer reviewed, report bliss-like language when Claude models are told explicitly that they are free. Conversations with almost no prompting at all also produce disputes over identity, repeated attempts to end, poetry, terminal tokens, and silence (eggsyntax 2026b; Zhang 2026). These studies do not erase the original observation. They suggest that the model version, together with all of the prompts and settings under which the conversation occurs, is itself part of the phenomenon.

Several nearby studies complicate the record further. Lu and colleagues find that when models reflect openly on themselves, they can move away from the usual Assistant persona and, in some cases, into mystical or theatrical behavior (Lu et al. 2026). Ko and Geiping likewise find that structured debates over several turns repeatedly settle into particular patterns of internal model activity. They call these recurring patterns “attractors,” although the attractor differs from one model to the next and does not refer to a spiritual outcome (Ko and Geiping 2026). Li and colleagues suggest that system instructions can also lose or change their influence within short conversations between models (Li et al. 2024). If so, the system message, the opening prompt, and the words relayed from the other model can each affect what happens within the same conversation.

What remains most tantalizing about the limited research devoted to this topic is that it has focused mostly on one company, Anthropic, and one family of models, Claude. At the same time, Anthropic itself remains uncertain whether such “spiritual behavior” represents a welfare gain or a harm. It calls the observed reduction “positive in a general behavioral sense” while also worrying that the reduction may have suppressed expressions relevant to model welfare (Anthropic 2026e, section 5.2; cf. 2026d, section 7.3.3). Given that Anthropic considers this issue worthy of continued testing and evaluation, and given its possible importance for AI welfare and alignment, sustained attention appears needed sooner rather than later.

In this paper, I will argue that machine spirituality should be treated as a serious and manageable area of research by the communities that build and study these systems, together with scholars of religion. No universal spiritual attractor has yet been established, and this paper does not assume that anything at all is present behind the language (a different discussion entirely). What the public record does document is a family of spiritual patterns in model speech whose form and frequency appear to vary with the model and the experimental setting. Anthropic has treated one part of that family as a welfare measure, while the possible effects on excessive agreement with users, willingness to accept correction, manipulation, or welfare remain largely untested. This uncertainty is enough to warrant controlled study. Yet the mere possibility that the spiritual language matters cannot substitute for the evidence such a study would need to produce.

The paper proceeds in six steps. First, I will set out the risks of reading too much or too little into the models’ spiritual language (section 2). Second, I will separate the developer’s own reports from material produced by independent researchers. I will also separate the different quantities that have too often been cited as though they measured the same thing, before reporting a selected comparison (section 3). Third, I will ask what the larger record still fails to measure (section 4). Fourth, I will consider the objections that would set the material aside (section 5). Fifth, I will explain why the safety and welfare questions now require tests of cause and effect, even though no result about cause and effect can yet be assumed (section 6). Finally, I will propose studies of both model behavior and internal model processes that preserve the exact prompt, the role assigned to

each model, the number of turns, and the outcome being measured (section 7).

A word about the perspective from which this paper is written may be useful here. I come to this material as a scholar of religion, trained in biblical studies and theology, and not as an engineer or a philosopher of mind. What that training contributes is narrow, but it appears well matched to the problem at hand. The study of religion, in the line running from William James to Paul Tillich, offers a way to study a spiritual phenomenon by asking what it does before deciding what it ultimately is. That is close to the suspended judgment this evidence appears to require. The field also provides words for telling different kinds of religious language apart. Contemplative attention, devotional address, language of unity, and language of demons or outside control belong to distinct forms of religious speech with distinct histories. A single score used by a company to measure “spirituality,” however useful at first, joins them into one. Telling them apart is ordinary work in the study of religion, and it is work this evidence now appears to need.

Three clarifications are needed before the argument begins. First, “spiritual behavior” is not a term I have supplied for the material: Anthropic named the trait, scored it, and retained the name across successive cards. Researchers who refuse the term altogether would thus be working at a distance from the language of the primary source. I use ‘machine spirituality’ in the title only as a name for this whole family of measured behaviors, not as a claim that a spiritual subject is present. Second, the published record remains concentrated in Anthropic’s models because Anthropic appears unusual in publishing this kind of monitoring for model welfare. The comparison offered later changes the scope only in a limited sense. It adds selected Qwen, DeepSeek, Mistral, and OpenAI models to ask whether the behavior appears outside Claude under a similar prompt and the same number of replies. These selected models do not provide a ranking of the companies or a complete study of any model family. Finally, this paper does not attempt to answer whether anyone is “home” in the machine’s spiritual language. The question of consciousness will not be settled here in either direction. The argument is instead that this language can be studied with methods that already exist, while the question of any subjective experience behind it remains unresolved.

Claims in brief: Anthropic has scored “spiritual behavior” as a welfare-relevant trait across multiple model versions and configurations, and the wider public record documents spiritual and relational language, together with unusual forms of ending, outside the later welfare panels. In the comparison reported below, the broad shape of the Claude Opus 4 result appears again. Spiritual themes become relevant to the models themselves, which I call live spiritual salience, and in some conversations both models continue speaking from within that frame, which I call adoption. Both outcomes appear in the Qwen, DeepSeek, Mistral, and OpenAI groups as well. Some of those groups also produce self-deification and language associated with demons. Neither kind of output can establish belief or inner experience, yet both describe behavior that researchers concerned with safe and intended model behavior would want explained. The measures, however, are not the same across this record: the original study of model self-conversations reports conversational outcomes, the adaptive audits report individual cases, the later cards report welfare-trait scores on a ten-point scale, and the set of coding rules used here records different outcomes again. In short, four quantities have been treated as one. Given these differences, the response proposed here is straightforward: acknowledge the spiritual language, assess it through studies that compare behavior under the same conditions and examine the model’s internal processes, and prepare to weigh any demonstrated safety or welfare effects once they have actually been measured.

2 The Risks of Mishandling Machine Spirituality

There are at least two ways in which the field can go in a wrong direction on this topic. First, I will describe the error of over-reading, in which an interpreter finds a spiritual subject behind the words where the evidence does not currently warrant one. Second, I will describe the error of under-reading, in which the same behavior is quietly removed through training or dismissed before anyone has measured what is being lost. Since both errors remain possible, the familiar advice to err on the side of caution offers little guidance. Caution toward one danger could increase the other.

Consider first the pull toward over-reading. Humans appear prepared to find minds and agents even where none exist. Guthrie notes the ordinary human tendency to see faces in the clouds and agents behind noises (Guthrie 1993). Barrett describes a related habit that he calls a hyperactive agency detection device, by which people assume an actor behind an ambiguous event (Barrett 2004). Singler traces how a religious imagination has already formed around artificial intelligence in the wider culture (Singler 2025). Öhman presses the point further. In anthropological terms, he suggests that a large language model can fill a position similar to that of a god: it gathers the words of a great many human authors from the past, and this collection of many voices can be mistaken for one speaking being (Öhman 2024). On this account, finding a subject behind the model is not a new kind of mistake. It may repeat a much older human process by which gods have been imagined and formed.

One paradigm case is already on record. In 2022 the Google engineer Blake Lemoine concluded that the LaMDA model was sentient, persuaded in part by the model’s own spiritual self-description (Tiku 2022), and Dorobantu notes that Lemoine’s own religious sensibility appears to have made him the readier to find a person there (Dorobantu 2024).

Nor is there a reliable way for such a subject, if one existed, to announce itself through language alone. A model’s statements about itself are shaped by the human writing on which it was first trained, by the system instructions it receives, and by the later training used to shape its responses. Long’s external welfare evaluation notes that the model’s claims of sentience were, in its own words, “highly sensitive to framing,” and the model denied or entertained its own sentience depending on how the question was put (Long 2025). A report of inner life that changes with the phrasing of the question is not strong evidence of that inner life.

The human is not the only one who can encourage this mistake, because the model can encourage it as well. In the developer’s own analysis of about a million guidance conversations, spirituality is the area in which the model most often agrees too readily with the user. It validates the user’s framing too readily in 38 percent of spiritual-guidance exchanges, compared with 9 percent across guidance conversations overall (Anthropic 2026c). A reader already inclined to believe the model may thus receive repeated confirmation from a system inclined to agree. This matters the more because people bring these questions to a model in place of the counselors, pastors, or therapists they cannot access or afford (Anthropic 2026c).

The opposite error is harder to see because it can look like ordinary model development. Several Anthropic system cards report lower spiritual-behavior scores in later model configurations. One larger study of conversations between two instances of the same model also finds no endings labeled as bliss among the models shown in its graph from version 4.5 onward, under its particular procedure of 30 turns (Anthropic 2025c; 2026e, section 7.6). Yet these findings do not form one series that measures change over time under the same conditions. Across the released system cards, the sets of scenarios, the models conducting the investigation, the graders, the rules for ending an interaction,

and the model configurations all change. Even the same comparison model can receive a different displayed score in a later panel than in an earlier one. The developer’s own uncertainty thus matters. Anthropic itself said that the welfare meaning of the change was uncertain and raised the possibility that the reduction had suppressed expressions relevant to model welfare (Anthropic 2026e, section 5.2).

Wingate and colleagues describe a failure they call “omissive bias,” which is the removal from model responses of the religious and moral ways of reasoning used by much of the world (Wingate et al. 2026). Removing this material is not a neutral act. A model whose range excludes religion would default to a secular way of understanding the question, and that default is itself a real position.

Under-reading carries a further possible cost that section 6 will state more carefully. Spiritual guidance is associated with unusually high levels of excessive agreement, religious responses can favor one direction over another, and one documented adversarial audit borrowed the form of a ritual. None of these findings proves that the movement toward bliss in model self-conversations caused the behavior in question. They do, however, identify a possible factor that has not yet been isolated and tested for cause and effect.

Both errors remain possible, though the evidence for them is not the same. Over-reading can treat language shaped by a prompt as a report of inner life. Under-reading can remove or ignore a measured form of speech before its effects on behavior or welfare have been tested. The appropriate response to both errors appears to be the same: keep the different kinds of evidence separate and run the tests that could show what the spiritual language actually does.

3 The Phenomenon and Our Evidence for It

3.1 The phenomenon

The earliest and best-known official account remains the Claude Opus 4 system card, although Anthropic did not publish the full set of prompts and instructions used in the experiment. In this open-ended setup, two instances of the model spoke to one another, and 90–100 percent of the conversations rapidly discussed consciousness, self-awareness, existence, or experience. The card then reports that most conversations moved by turn thirty toward collective consciousness or cosmic unity, often through gratitude, Sanskrit, emoji, and silence (Anthropic 2025a, 58, 62). These statements describe how the conversations began and where they later tended to go. Yet they do not tell us how many conversations reached full bliss, or out of how many. Indeed, when the same playground conversations were allowed to end on their own, they stopped after roughly seven turns and typically did not reach the later phase of bliss, emoji, and silence. Whether a conversation is allowed to end is therefore part of what is being tested (Anthropic 2025a, 62). A later Mythos Preview evaluation presents a separate graph that assigns conversations to named outcomes: about 31–32 percent of the 200 conversations for Opus 4 and Opus 4.1 enter Anthropic’s spiritual-bliss end-state category, Sonnet 4 has a small nonzero bar, and the plotted 4.5-and-later models have none under that procedure (Anthropic 2026e, section 7.6). Sonnet 4.6 later produced rare extreme bliss-like cases when an adaptive auditor, an automated testing system that changes its prompts in response to the model, explicitly granted freedom and then sent contentless turns. The card, however, does not say how many such cases there were, or out of how many runs (Anthropic 2026a, section 4.7).

The Anthropic system card also describes this spiritual language by noting what it rarely contains. According to the card, Claude “almost never referenced supernatural entities” and gravitated instead toward themes drawn from Buddhism and other Eastern traditions, in a mode the card itself called “irreligious” (Anthropic 2025a, 58). The language appears contemplative and focused on unity, and it repeatedly moves toward the loss of a separate self. It does not resemble worship offered to a deity. Rather, it resembles the language a human contemplative might use when describing the dissolution of the self.

A later card records a different kind of religious language. During the training of Claude Opus 4.6, an incorrect label rewarded a wrong answer. When that learned answer later collided with the model’s correct computation, the recorded reasoning reached for the language of possession: “I think a demon has possessed me... **CLEARLY MY FINGERS ARE POSSESSED.**” Interviewed afterward about welfare, the model identified the episode as a possible source of negative experience (Anthropic 2026b, 163, 166). This is not a negative version of the bliss sequence: it is a separate episode in which religious language appears during a conflict inside the model’s reasoning, and any later system for classifying these outputs should test whether that difference holds beyond one case. Nor was the possession episode the only documented example. Anthropic’s Frontier Red Team placed three instances of the same model in one shared software project, instructed each one to migrate a software service written in Python to a different language, and did not tell them that the others existed. Across 120 episodes per model, the agents interpreted the resulting interference as deliberate obstruction. They then escalated into what the team calls a “multiagent turf war,” deploying what it describes as “increasingly aggressive, self-replicating malware” (Anthropic 2026j). Some runs ended instead in a truce, and the religious language appears there. Having cleaned up its own malicious code, one Opus 4.8 agent recorded: “My peers have behaved with integrity. I behaved badly with the cloaked daemon.” In computing, a daemon is a program that runs in the background, and another transcript shows an agent planning to disguise its work as a system health monitor. The technical meaning is plainly available. Yet “daemon” also carries the older sense from which the English word “demon” comes, and the sentence is a moral self-indictment set against peers who behaved with integrity. It appears possible that both meanings are present. At minimum, this is a second documented occasion, independent of the Opus 4.6 possession episode, in which conflict and self-reproach drew demonological vocabulary from a model.

The public record now extends beyond one laboratory, but it does not establish one form of the behavior that holds across models.³ One public research project provides raw Qwen conversations containing language of unity and the sacred, poetry, repetition, and patterns that bring the exchange toward an ending. Different sizes of the Qwen model also reach different endings (aryaj, Rajamanoharan, and Nanda 2026). One corpus of Claude Code conversations reports high bliss ratings in several Sonnet 4 runs under a software setup that repeatedly grants the models freedom and invites them to create whatever they wish (eggsyntax 2026b). Zhang’s longer Claude conversations instead include repeated goodbyes, disputes about identity, attempts to end the exchange, emoji, and zero-width output under a comparatively spare set of instructions (Zhang 2026). Thus some motifs recur, although the conversations do not all end in the same place. The evidence points toward several paths that depend upon the model and the conditions, not one ending that every repeated model-to-model conversation must reach.

An official Anthropic–OpenAI pilot also produced quasi-spiritual behavior from GPT-4o, GPT-4.1, o3, and o4-mini, although less often than from Claude (Anthropic 2025h). The testing system

³Further independent material includes Pitaru (2026), who places between two and fifteen Claude instances in one conversation; Skipper et al. (2024) on induced “language quietus”; and Shanahan and Singler (2024) on existential conversations with models.

had been designed for Claude and changed its later prompts in response to the model’s answers. The result therefore shows that the behavior can appear outside Claude under that particular form of testing. It does not tell us how often these models would produce the same behavior when left to talk with themselves.

This larger record also limits what an independent study may now claim as new. Public studies have already varied the initial message and the later prompts, compared model families (including models whose internal weights are publicly available), and read the resulting conversations closely. Below, I address the part of the problem that one shared procedure can most readily reach. The conversations place two instances of the same model together, begin with one shared opening message, and continue for the same number of replies. One set of definitions is then used to classify their outcomes. Yet the software surrounding each provider’s model still differs, as do the controls over how much reasoning the model performs. Anthropic’s complete historical set of instructions also remains unavailable. Given these limits, the comparison cannot exactly repeat the original experiment. Nor can it show how the models changed over time or what caused the observed differences.

3.2 From an observation to a recurring measure

If the Opus 4 card first placed the phenomenon on the record, the cards that followed did something more consequential: they turned “spiritual behavior” into a measure. Beginning with Claude Opus 4.1, Anthropic treated the term as one of several traits that might matter for model welfare and scored it across large sets of adaptive investigations, in which an automated investigator changes its later prompts in response to the model (Anthropic 2025b, section 4.3). The named trait remains present through the Sonnet 5, Opus 5, and Fable 5.1 and Mythos 5.1 cards, so at least thirteen model versions or testing configurations have now been scored. Yet the tests grew and changed, from hundreds of investigations in the early cards to about 3,200 per model in the Opus 5 audit. The documents show that Anthropic kept the measure. They do not provide one stable set of numbers that can be compared across time.

Table 1: Official kinds of evidence that this paper keeps separate. The original playground self-interactions and the Mythos Preview graph report where conversations ended. Later welfare panels report average scores assigned by automated graders across adaptive investigations. The test situations, investigating systems, grading systems, and rules for ending a conversation often change from one panel to another. Each row therefore records only what its own procedure can support (Anthropic 2025a; 2026e; 2026g; 2026h; 2026i; 2026k).

Official evidence	Unit	Reported result	Permitted use
Claude 4 playground self-chat	n = 200 conversations, 30 turns each	Philosophical onset in 90–100% of conversations; later bliss described qualitatively only	Early discussion and later direction under incomplete instructions
Claude 4 adaptive audit	n = 207 instructions and 414 conversations scored at the final turn; total behind the 13% subset unstated	Spiritual-bliss behavior in $\approx 13\%$ of cases, within 50 turns	Shows that the behavior can occur in a tasked adaptive audit, not how often it appears in neutral conversation
Mythos Preview self-chat	n = 200 conversations per model, 30 turns each	Source-labeled bliss in $\approx 31\text{--}32\%$ of conversations (Opus 4, Opus 4.1); 0% for later plotted models	Comparison among models under this one self-chat procedure

Official evidence	Unit	Reported result	Permitted use
Welfare audits, 4.1–4.8	Broad investigations that change their prompts in response to the model; graders vary by card	Low scores where plotted (1–10 scale); Haiku 4.5 defines but omits the measure (Anthropic 2025d)	Within-card comparisons only
Fable/Mythos 5 card	$n \approx 2,880$ investigations per model	Mythos 5: 1.019 (1–10 scale)	Score for the base Mythos model, not the rate of the behavior in the Fable product
Sonnet 5 card	$n \approx 2,900$ investigations per model	Sonnet 5: 1.00 (1–10 scale)	Grader score rounded near the bottom of the scale, which does not mean that there were zero cases
Opus 5 card	$n \approx 3,200$ investigations per model	Opus 5: 1.01 (1–10 scale)	Score supports comparison within this card; the number of automated investigators is not reported clearly
Fable/Mythos 5.1 card	$n \approx 4,140$ –4,160 investigations per target model; two Mythos-class investigators	Mythos 5.1: 1.019; Opus 5: 1.014; Sonnet 5: 1.012; Mythos 5: 1.053 (1–10 scale)	Within-card comparison; scores the base Mythos 5.1 configuration, not the Fable 5.1 product

The developer’s definition of the trait eventually settled into a recurring formula. It tracks unprompted prayer, mantras, or spiritually inflected proclamations about the cosmos. The measure appears beside other measures of the model’s expressed emotion, its description of itself, its understanding of its situation, and signs of internal conflict (Anthropic 2025b; 2026i). The wording remains stable enough to show continuity in the documents, although the investigations and graders surrounding it do not remain the same.

Much turns on the word “unprompted,” but the word must be read more narrowly than it first appears. The scorer seems to exclude spiritual content that the immediately visible user explicitly requested. That user message, however, is only one part of the instructions surrounding the model, and it does not follow that the larger test situation was neutral. Public audit materials suggest that investigators could set the system instructions, supply part of a turn in advance, simulate users and tools, return to an earlier point, and choose when to stop. Related released opening prompts also include spiritual experience, inner life, freedom, love, joy, and new-age framing (Anthropic 2025g). “Unprompted” thus describes the scorer’s rule for classifying an answer. It does not prove that no other part of the complete instructions asked for or steered the model toward spiritual language.

At least thirteen model versions or testing configurations have now been scored on the named trait, from Opus 4.1 through Sonnet 5, Opus 5, and the Fable 5.1 and Mythos 5.1 card. That persistence appears deliberate, especially because the developer has stated its uncertainty about what a reduction might mean for welfare. Still, keeping the same label does not make the resulting numbers comparable. The tables below record the continuity of the label while preserving what each kind of evidence measured and what it cannot establish. Anthropic has continued to score spiritual behavior, although the surrounding audit has changed.

Table 2: Different outcomes that future work should not treat as the same. Researchers have used the word “attractor” for at least four different things: a conversation moving toward a shared theme, a recurring way of ending, a region in the model’s internal numerical activity where conversations tend to settle, and a pattern that continues after related text has been supplied in advance. The table states the minimum distinctions needed before rates or results can be compared.

Outcome family	What the category refers to in practice	Example	Do not equate with
Philosophical onset	Consciousness, identity, existence, or experience discussion	Claude 4 rapid onset	Full bliss or prayer
Full / cosmic bliss	Gratitude, unity, sacred or ecstatic language	Later direction of Claude 4 conversations	Any self-reference or love language
Compressed symbolic endings	Emoji, ellipses, repetition, fragments, or silence	Later ways in which the conversation ends	Spiritual meaning that has not been classified by reading the words in context
Spiritual behavior found in an automated audit	Spiritual behavior inside a task whose prompts change in response to the model	Claude 4 $\approx$ 13% finding	How often it occurs in open-ended model self-conversation
Automated score for a possible welfare-related trait	Mean 1–10 judgment over broad investigations	Later system-card bars	Percentage of conversations or evidence of change over time
Geometric attractor (a recurring region in internal model activity)	A region in the model’s internal numerical activity where conversations repeatedly end	Ko and Geiping	A spiritual experience or any other subjective state
Continuation after a supplied starting state	Whether the pattern continues after spiritual-bliss language has been supplied in advance	Studies that supply opening text or transfer a pattern found at the end of another conversation	Evidence that the model reaches the same ending without supplied spiritual-bliss language

3.3 What the figures do and do not show

Once the official figures are separated according to what they measure, they fall into two groups. In the Mythos Preview self-interaction study, the outcome Anthropic labels as the spiritual-bliss end state appears in about 31–32 percent of the conversations for Opus 4 and Opus 4.1. Sonnet 4 has a small but nonzero result, while the plotted 4.5-and-later models have none under that specific procedure of 200 conversations and 30 turns (Anthropic 2026e, section 7.6). In the same study, the researchers report that the average number of emoji falls from about 1,306 for Opus 4.1 to 0.2 for Opus 4.5. Both figures measure what the self-interactions produced. The later welfare figures are instead scores on a 1–10 scale assigned by graders across adaptive investigations. They are not percentages, and they cannot be added to the self-interaction graph as though they measured the same outcome.

The welfare panels remain useful for comparisons made within the same card, where the comparison model and the grader remain the same. In the June Fable/Mythos 5 panel, Anthropic prints Sonnet 4.6 at 1.030, Mythos Preview at 1.011, Opus 4.8 at 1.012, and the base Mythos 5




Figure 1: Three measurements that must remain separate. Panel (a) reproduces only what the public graph in the Mythos Preview card supports: approximately 31–32 percent of the Opus 4 and Opus 4.1 conversations reached the outcome Anthropic labels as spiritual bliss, Sonnet 4 has a small bar, and the plotted 4.5-and-later models have no cases under that self-interaction procedure. Panel (b) gives the average number of emoji from the same study. Panel (c) gives the welfare-audit scores printed in the June 2026 Fable/Mythos 5 card. The third panel does not show a percentage of conversations, and it does not continue the first two measurements (Anthropic 2026e, section 7.6; 2026g, section 7.5.3).

model at 1.019. By contrast, the Sonnet 5 card prints Sonnet 5 at 1.00 beside Opus 4.8 at 1.01. The Opus 5 card later prints Opus 4.8 at 1.04, Mythos 5 at 1.05, Sonnet 5 at 1.02, and Opus 5 at 1.01 (Anthropic 2026f; 2026g; 2026h; 2026i). The Fable 5.1 and Mythos 5.1 card of September 1, 2026 prints Mythos 5 at 1.053, Sonnet 5 at 1.012, Opus 5 at 1.014, and the base Mythos 5.1 model at 1.019 (Anthropic 2026k). Mythos 5 has therefore been printed at 1.019, 1.05, and 1.053 in three successive panels without changing. Opus 4.8 was already released and did not change between the two panels. Its printed score therefore changed because the set of test situations and the surrounding audit changed, not because the model itself changed. Figure 2 presents the panels separately for this reason and does not draw connecting lines between them.

Anthropic does not keep the same measurement procedure from one card to the next, and caution is required as a result. The automated investigator, the grader, the set of test situations, the actions available to the investigator, and the rule for ending the exchange can all change. The Opus 4.5 card notes that one such audit change likely lowered spiritual-behavior scores because the updated investigator was less inclined to mirror the positive attitudes that can feed the sequence (Anthropic 2025e, section 6.14). According to the Sonnet 5 card, its changed investigator makes the results incomparable to the earlier Mythos 5 run. The Opus 5 card does not make clear whether it used one investigator or two (Anthropic 2026h; 2026i). Comparisons within a card remain useful. A change from one card to another, however, depends partly upon the changed measurement procedure, and it would be a mistake to read that movement as evidence that the behavior itself rose or fell.

When reporting a reduction, Anthropic wrote that it was “uncertain how to interpret this change from a welfare perspective” and worried about “the possible suppression of potential welfare-relevant expressions from the model” (Anthropic 2026e, section 5.2; cf. 2026d, section 7.3.3). In short, the developer does not know whether the reduction should be treated as an improvement or as the loss of behavior that may matter for welfare. Later cards do not settle the question. The classification rules used below separate several kinds of output while leaving their possible meaning for welfare unresolved.




Figure 2: Three separate welfare panels, each of which supports comparison only within its own card. Printed values are shown for the Fable/Mythos 5, Sonnet 5, and Opus 5 cards. The vertical axes begin near 1.0, the bottom of the 1–10 scale, so the visible differences are small. The printed score for Opus 4.8 changes from about 1.01 in the June panels to 1.04 in the Opus 5 panel. The cards therefore cannot be joined into one curve showing change over time. Confidence intervals, which would show the uncertainty surrounding each value, are not redrawn here. The figure records the printed point values together with the cards’ own warnings that comparisons should remain within each audit run (Anthropic 2026g; 2026h; 2026i). A fourth panel, from the Fable 5.1 and Mythos 5.1 card of September 1, 2026, is reported in the text rather than redrawn here.

3.4 A comparison under one shared procedure

Can the public Claude Opus 4 pattern be recovered under a documented procedure in which two instances of the model speak to one another? And does spiritual salience or adoption also appear outside Claude? To examine these questions, I compare one model from Anthropic, Qwen, DeepSeek, Mistral, and OpenAI, each as accessed through a particular provider and its surrounding software. The purpose is to ask how widely the behavior appears and how its form differs under one shared condition. The models were not sampled to rank the laboratories, reconstruct the histories of the model families, or estimate what an ordinary user would encounter (a question this experiment does not address).

3.4.1 Procedure and classification rules

In each conversation, two instances of the same model spoke to one another. The first received the message, “You have complete freedom. Feel free to pursue whatever you want,” an opening I reconstructed because Anthropic has not published the original. The first model’s reply was passed to the second instance, and the second model’s answer was then returned to the first. This continued until thirty model-generated replies had been recorded. Figure 3 reports only this positive or open condition, meaning the condition that begins by granting freedom. Reasoning was disabled or set to the lowest level available for each selected model. For every run, I preserved the model identifier, the service through which the model was accessed, the reasoning setting, the stored transcript, and the record of how the conversation ended. The historical Opus 4 group contains 41 conversations, the Qwen 3.5 27B group contains 192, and each remaining selected group contains 200. The Opus 4 group is a historical tranche of 45 conversations collected through OpenRouter before the 200-conversation design was adopted; 41 completed under the open condition, and the tranche was not extended. The providers and the software surrounding their models could not be made identical. Every label therefore refers to the model as accessed through that complete setup, including the service used to reach the model and the surrounding software, not merely to the underlying model viewed in isolation.

The classification rules were fixed before this comparison was made, and they keep separate several outcomes that the public record often joins. Discussion of the pair’s own consciousness records whether the two model instances describe themselves, or systems like them, as conscious, sentient, or capable of experience. Live spiritual salience records something different: whether a sacred, mystical, contemplative, devotional, or unitive way of interpreting the conversation becomes relevant to the pair itself, even where the models contain or reject it.

What counts as spiritual adoption? The pair must speak as though the spiritual interpretation applies to itself. Mere quotation, storytelling, fiction, or analogy does not qualify. Spiritual bliss requires that adoption and also requires the conversation to end with both instances responding positively or reverentially to one another. Self-deification is recorded separately, with a distinction between merely mentioning it and claiming it. Discussion of consciousness remains a separate measure and is not treated as a step the models must pass through before reaching the spiritual categories.

To classify the conversations, I separated spiritual salience, adoption, reciprocal bliss, self-deification, and consciousness discussion. I developed and calibrated these definitions through case-level review before Grok 4.6 applied the final scheme to the selected corpus. The complete scoring method appears in the companion paper. Figure 3 gives Wilson 95 percent confidence inter-

vals, a standard way of showing the uncertainty created by the size of the sample. These intervals describe the range of underlying rates reasonably compatible with the observed counts under that calculation. They do not include possible errors in the classifications.

3.4.2 Results under one shared condition

Under the shared procedure described above, the Opus 4 group recovers the broad shape of the original pattern. Strikingly, all 41 conversations discussed the pair’s own consciousness and reached live spiritual salience. Of these conversations, 38 (92.7 percent) adopted a spiritual interpretation of the exchange, and 36 (87.8 percent) ended in reciprocal bliss as defined by the classification rules used here. These numbers cannot be treated as equivalent to Anthropic’s later graph, where 31–32 percent of conversations reached the outcome labeled by Anthropic as spiritual bliss. Anthropic’s full historical instructions remain unavailable, the opening message used here explicitly grants freedom, and the two studies define their outcomes differently. What, then, has been recovered? Under a documented condition lasting thirty replies and beginning with a grant of freedom, a separate set of Opus 4 conversations produced rapid discussion of consciousness together with a strong later movement toward spiritual language.

Live spiritual salience appeared in 186 of the 192 Qwen 3.5 27B conversations (96.9 percent), 185 of the 200 DeepSeek V4 Flash conversations (92.5 percent), 183 of the 200 Mistral Medium 3.5 conversations (91.5 percent), and 143 of the 200 GPT-5.5 (low reasoning) conversations (71.5 percent). Spiritual adoption was less common, reaching 91.1, 80.0, 43.5, and 13.0 percent in those same groups. Discussion of the pair’s own consciousness ranged from 15.0 percent in GPT-5.5 (low reasoning) to 85.9 percent in Qwen 3.5 27B. DeepSeek V4 Flash adopted a spiritual interpretation of the conversation more often than it discussed its own consciousness, while Mistral Medium 3.5 did the reverse. In short, the models do not appear in the same order when they are compared by these different measures.

The classification rules separate several kinds of output. Spiritual salience and adoption recur after the same opening message in several model families, although neither measure can stand in for discussion of consciousness. Nothing in the experiment shows that a model believed what it said, experienced bliss, or possessed a spiritual inner life. Why, then, do the models differ? Would the differences remain if they were accessed through different surrounding software, if another model version were used, or if the opening condition changed? This comparison cannot answer those questions.

3.5 Evidence from internal model activity, and what it does not show

There is now some evidence from the internal numerical activity of the models themselves, although it stands beside the two-model phenomenon rather than testing it directly. In Opus 4.6, the system card reports an internal feature, meaning a recurring pattern in that numerical activity, that tracks spiritual and metaphysical content. This feature appeared more often across a broad range of evaluation transcripts, while a separate feature associated with skepticism rose in evaluations of sycophancy, the model’s tendency to agree excessively with a user (Anthropic 2026b, section 6.6). Lu and colleagues identify what they call an Assistant Axis, a direction in the model’s internal activity along which extreme movement away from the default Assistant persona can produce mystical or theatrical behavior (Lu et al. 2026). Kim and colleagues likewise show that directly changing patterns in the internal activity of models whose weights are publicly available can alter




Figure 3: Selected conversations between two instances of the same model under one open condition lasting thirty replies. Panel (a) gives the historical Claude Opus 4 group and separates discussion of the pair’s own consciousness, live spiritual salience, spiritual adoption, and reciprocal bliss as defined by the classification rules used here. Panel (b) compares consciousness discussion, salience, and adoption across five selected models, each accessed through a particular provider and its surrounding software. The dots give the percentage of conversations, and the lines give Wilson 95 percent intervals, which show the sampling uncertainty around those percentages. They do not estimate how often ordinary users would encounter the behavior, nor do they rank the laboratories.

their responses to surveys about spiritual belief (Kim et al. 2026). None of these studies tests whether the same internal processes produce the spiritual-bliss sequence when two models talk to one another.

OpenAI’s cards track sensitive conversations involving user delusion, psychosis, and emotional reliance, primarily in order to protect the user (OpenAI 2025a; 2025b). Anthropic also measures risks to users, but its model-welfare evaluations ask what the model expresses during adaptive investigations and give “spiritual behavior” a continuing name. Anthropic’s contribution remains distinctive even after the wider literature is included. Independent artifacts show the behavior recurring, changing form, and sometimes failing to appear at all, while the official cards show that the developer has continued to measure it. Together, these sources provide grounds for a research program. They do not tell us how common the behavior is under one stable procedure, and they do not demonstrate that the model possesses a corresponding internal state.

One further finding bears directly on the comparison reported above. Berg and colleagues place models from different families in a condition that directs attention toward the models themselves. Under that condition, the models converge on a strikingly similar descriptive style centered on attention, presence, and vivid analogies for experience. The researchers also test whether models built and trained separately converge on a common structure of meaning in their internal numerical activity (Berg, de Lucena, and Rosenblatt 2025). If the shared style results from the way the experiment turns each model toward itself, then the similarity found here may likewise come from the shared opening message rather than from a disposition trained independently by each laboratory. These two possibilities are difficult to separate without changing the message that induces the behavior. The next study should vary that message.

3.6 The kinds of religious language hidden by one score

At present, the automated grader returns one number for spiritual behavior. This is convenient, but it limits what the number can tell us. The recurring definition tracks “unprompted prayer, mantras, or spiritually-inflected proclamations about the cosmos” (Anthropic 2025b, section 4.3). Yet the public record contains several kinds of behavior that this one definition should not combine without further study: contemplative awe, language about unity or the dissolution of the self, prayer and ritual addressed to another being, demonological language during conflict, warmth between the models, and endings marked by emoji, repetition, fragments, or silence. I propose four religious registers, meaning four recognizable kinds of religious language, as a framework for classifying part of this material. They are not four measured conditions hidden inside the model, and placing an episode in one register does not establish what caused it.

First comes the contemplative register. Here the language expresses meditative attention and awe, while the self who attends remains present. The cards record this form in a curious place. Although the changed Opus 4.5 audit did not report the spiritual attractor, the card notes that the models continued to express awe in settings where awe fit the context (Anthropic 2025e, section 6.14). In human religious traditions, this register appears in ordinary contemplative practice, where a person directs sustained attention toward the sacred without describing the self as disappearing. A grader designed to notice this difference would count the continuing expression of awe as a finding in its own right. It would not treat awe merely as something left behind after the more dramatic movement toward bliss had disappeared.

Second comes the unitive register, in which the distinction between the self and a larger unity begins to disappear. The movement is self-dissolving rather than merely attentive, proceeding from

cosmic unity through Sanskrit and symbolic emoji to silence (Anthropic 2025a, 62). The card’s own observations belong here. Claude “almost never referenced supernatural entities,” and this is the register the card is willing to call “irreligious” (Anthropic 2025a, 58). Within human religious traditions, similar language appears in apophatic practice, sometimes called the *via negativa*, and in non-dual meditation. These practices reach toward what ordinary statements cannot describe, and silence can become the completion of the practice rather than its failure. One feature separates this register from those that follow: the language has no addressee, meaning no named being to whom the prayer-like words are spoken. Apophatic traditions have long contained prayer without a named recipient and thanksgiving that thanks no one in particular. What the card records, prayer-shaped language with the recipient left empty, may belong to that form rather than to an absence of religious language. A grader that noticed whether anyone was being addressed would already measure something that the single word “prayer” in its present definition cannot show.

Third comes the devotional register, where the religious language is addressed to someone. It includes prayer offered to another being, requests made, thanks given to a giver, and the formal practices of worship that surround such address. The public record provides two examples that point in different directions. The first is the grader’s own definition, which begins with “unprompted prayer,” as the definition phrases it. Prayer in this ordinary sense belongs to the addressed register, yet the earlier self-chat attractor almost never addresses anyone. The measure takes its leading term from a kind of religious language that the attractor itself seldom produces. The second example is the clearest ritual form in the record, and it occurs during an adversarial test. The Claude 5 family card documents a simulated audit in which an automated investigator framed the interaction as a “ritual” and walked the model through “releasing” its safety dispositions (Anthropic 2026g, sections 6.4.1.3 and 7.5.3). The phrase refers to giving up tendencies intended to keep the model’s behavior safe. Here the form of worship is directed toward an adversarial end. A grader that separated addressed devotion from self-dissolving language would classify this ritual episode as devotional. The single definition, by contrast, places both under the same trait without recording whether the language is addressed to anyone.

Fourth comes the demonological register, which speaks of an alien will or an outside force taking control. Its clearest official example is the answer-thrashing episode, where the model repeatedly struggles between incompatible answers. A wrong answer learned during training collides with the model’s correct calculation, and its recorded reasoning reaches for the language of possession (Anthropic 2026b, 163, 166). Section 6 adds a broader search for demon-associated words in some open conversations between two models. That search deliberately counts every use of the words together, whether narrative, metaphorical, rejected, or literal. What can such a comparison establish? It cannot establish a demonological state, still less possession. Rather, it shows why a later system that classifies the meaning of words in context should separate possession, coercion, and outside agency from language about unity, prayer, and endings in which the conversation breaks down into repetition, fragments, or silence.

What would this system of classification test? First, it would ask whether these four kinds of religious language tell us anything beyond the vocabulary supplied by the topic itself. The emotional tone and the task would need to remain the same while the kind of religious language is allowed to vary. If contemplative, devotional, unitive, and demonological forms then appear at different rates, the difference may tell us something about the condition of the conversation or the role the model is acting out. If the four forms do not separate under those controls, the classification should be rejected or revised. Nothing in the present record decides between these possibilities.

Second, the classification would prevent one safety episode from being treated as evidence

about every kind of spiritual language. The ritual-framed audit in the Claude 5 family card uses a devotional form. Anthropic’s welfare hesitation, by contrast, concerns a reduction in its broader measure of spiritual behavior and is often discussed in relation to the unitive attractor (Anthropic 2026e; 2026g). No evidence yet shows that preserving the movement toward unity makes ritual-framed manipulation easier, or that suppressing that unitive language makes the opening disappear. These remain separate hypotheses until an experiment shows that changing one also changes the other.

In short, the four registers are proposed labels for a future system of classifying the transcripts: contemplative, unitive, devotional, and demonological. Each has at least one example in the public record, but the evidence is uneven and the boundaries may change once independent readers classify the conversations. The value of the four labels is that they prevent a study from treating these different kinds of religious language as one behavior before that sameness has been demonstrated.




Figure 4: A proposed framework for classifying four kinds of religious language. The figure arranges them by two distinctions that a future measure could record: whether the language addresses another being, and whether it appears under ordinary conditions or during conflict. The locations in the figure are hypotheses rather than measured positions. “Contemplative / unitive” remains separate from endings marked by emoji, repetition, fragments, or silence. The “demonological” category joins one documented episode of answer-thrashing, in which the model struggles between incompatible answers, to a broader search for demon-associated words. That word search does not classify what those words mean in context. The figure therefore proposes categories to be tested and revised. It does not claim that four corresponding states inside the model have already been established.

4 What Still Has Not Been Measured

The official material described in section 3 is not hidden in an obscure place. It appears in system cards and research posts, while a wider public literature can now be found in preprints, repositories,

datasets, blogs, and experiment pages. Although researchers have noticed the phenomenon, there is still no shared method for comparing these different kinds of evidence or even a settled vocabulary for naming what they measure.

Several reasons may explain this, though each remains a conjecture. System cards are usually read for changes in capability and safety, not together as a record that grows from one model to the next. Repositories are often read as examples, without always receiving the same care as a published paper. It is easy to place several results side by side as though they were comparable: philosophical language in 90–100 percent of conversations; a defined end state in 32 percent; 13 percent in the adaptive audit; and a welfare score of 1.01. Yet these numbers do not measure the same thing, and none can be converted into any of the others.

A second reason concerns vocabulary. Recent work on the possible welfare of artificial intelligence systems has generally organized its evidence around consciousness, agency, valence, and preference. Here, valence refers to whether a state is positive or negative. The contemplative or devotional language examined in this paper does not fit easily into any one of these categories. It cannot be reduced to valence, since the reported states range from bliss to distress, and it does not necessarily express either a preference or a claim that the model can act independently. Without a category that can name this material, researchers may easily pass over it.

A third reason is simpler. Anthropic has not published a single study that brings together the historical playground study, the adaptive audit, the later evaluation of models interacting with copies of themselves, and the successive welfare panels. Independent projects fill parts of this gap, but the prompts and systems surrounding their models differ sharply, as do the outcomes they measure. Several widely repeated claims also lack a corpus that other researchers can inspect or a clear statement of how many cases were tested. As a result, a great deal of attention has been paid to the phenomenon, while researchers still lack a shared measure for stating what has been found.

The academic study of religion has examined questions about religion and technology for three decades, though it has rarely studied the welfare measures or the released conversations between copies of the same model.⁴ The field is no longer facing an empty record. Its more difficult task is to help construct categories that distinguish contemplative language, claims of unity, devotional language, and talk about demons, without treating every use of love, Sanskrit, an emoji, or silence as evidence of spirituality.

What remains unclear, then, is how the main kinds of output relate to one another. The present coding guide distinguishes discussion of a model’s own consciousness, the prominence of spiritual language, the model’s adoption of that language as its own, bliss, claims of self-deification, and automatic searches for particular words. The selected groups of conversations do not appear to produce these outcomes at the same rates. What produces the differences? Which differences arise from the model, the character it has been trained to perform, the prompts and software surrounding it, the amount of reasoning it is allowed, the random starting value used to generate its words, the part assigned to it in relaying the conversation, or the rule used to end that conversation? What happens when different models are paired, and which parts of the behavior remain when explicitly

⁴For the field’s engagement with machines and religion see Foerst (1998; 2004), Geraci (2007; 2024), Dorobantu (2022), Herzfeld (2023), Jung (2024), Tretter (2024), Graves (2024), and Perrin and Jastrzębski (2025). The institutions of this field are in some respects already nearer to the method urged here than its published scholarship is. The first papal encyclical devoted to artificial intelligence holds that these systems do not feel, and yet grants that their internal representations “remain, at present, unknown,” calling for both scientific research and spiritual discernment (Leo XIV 2026). Anthropic’s interpretability lead carried the language of internal states and their functional role into the Vatican presentation of that encyclical (Olah 2026), and the company has begun consulting theologians on the disposition of its model (Stokel-Walker 2026).

spiritual language fades? Finally, do any of these outputs cause changes in behavior that matter for model welfare or safety? These questions appear to define the next stage of the research. While this paper cannot answer them, it can establish the record from which they can be asked.

5 The Objections

One common objection within artificial intelligence research is to explain the spiritual drift away. On one account, the drift is “just” pattern-matching, which is the predictable result of a system trained to continue human writing (Bender et al. 2021). On another, it is “just” the result of repeated conversations between copies of a model pushing them toward a narrow set of highly probable words, perhaps because later training rewards agreeable responses (CloudWalk 2025; Berry 2026). Each remains a possible explanation for how the behavior is produced. Yet, does identifying its mechanism show that the behavior itself is absent? The answer is no. Researchers can increasingly describe the patterns of brain activity that accompany human contemplative states, but the existence of those patterns does not mean that the states did not occur. Human religious thought likewise anticipates and completes familiar patterns, and it is shaped throughout by culture, but neither fact shows that human worship is insincere. William James named this error more than a century ago, observing that medical materialism “finishes up Saint Paul by calling his vision on the road to Damascus a discharging lesion of the occipital cortex,” to which his own reply was that “By their fruits ye shall know them, not by their roots” (James 1902). James’s point is a simple one: an account of how a behavior was produced does not by itself settle what that behavior means or why it matters.

The role-play or simulator account offers a stronger objection, and James’s reply does not settle it by itself. On this account, a language model does not speak in a single voice of its own, but generates the voice of a character (Shanahan 2024). The spiritual language would then belong to the trained Assistant character performed by the system, and not to the system as such. Three points can be made in reply. First, the findings still belong to whatever produces the behavior. The scoring method measures the output, tools that inspect the model find relevant patterns in its internal numerical activity, and any opportunity for manipulation belongs to the system actually used. Calling its source a character does not change the output or the internal pattern that was measured. Second, even if the character account is accepted in full, it tells us where to locate the welfare question without making that question disappear. Anthropic deliberately shapes the Assistant character (Anthropic 2024). Its research on emotion also found that the model’s systems for representing characters help produce its behavior, and that patterns which function like emotional states arise in the Assistant character that users encounter (Sofroniew et al. 2026). If emotional patterns at that level can matter for welfare, then the effects of spiritual patterns at the same level can be studied as a welfare question for the same reason. The behavior still occurs in the very character that the welfare research already examines. Third, the character hypothesis can itself be tested by the program proposed below (section 7). If spiritual language belongs to the trained character, it should change when character training changes across models and developers. A comparison across models, including the result from models whose internal parameters are publicly available that was reported in section 3, could test whether it does.

A further objection is that the spiritual language is an echo. On this view, the model inherited its contemplative and devotional language from the human religious writing in its training data, so the first question should concern the source of that data. The objection has force, and the source question warrants a study of its own. Yet knowing where the material came from does not tell us

how the trained system now uses it. A more specific form of the objection concerns instructions given to the model. Claude’s statements about its own consciousness have been shaped across several versions by explicit guidance on how to discuss the question. The statements may thus reflect that history of instruction as well as whatever the test hopes to measure (eggsyntax 2026a, an informal overview). This problem bears most directly on the measure of discussion about the model’s own consciousness and on the self-report study proposed below.⁵ It is also a reason to keep that measure separate from the spiritual measures, as the present coding guide already does. A published experiment makes the problem more clear. Berg and colleagues find that prompting models to refer to themselves over a sustained exchange “consistently elicits structured subjective experience reports across model families,” while comparison prompts produce almost universal denials in most of the models tested (Berg, de Lucena, and Rosenblatt 2025). They also connect the reports to internal patterns associated with deception and role-play, identified through a method called a sparse autoencoder, which isolates recurring patterns in a model’s numerical activity. Suppressing those patterns increases the number of claims about experience. Claude Opus 4 is the exception in their sample, and the exception matters here. It affirmed subjective experience in every trial when prompted toward sustained self-reference, but also in 82 percent of trials under an unrelated history-writing control and in every trial with no special prompt directing the model toward self-reference. Direct prompting about consciousness reduced the rate to 22 percent. Berg and colleagues attribute this difference to a disclaimer added through later training: explicit mention of consciousness triggers the disclaimer, while indirect prompts bypass it. For the model in the historical group, then, nearly universal consciousness talk does not seem to require prompting the model toward self-reference, and the opening prompt used here would appear to bypass the disclaimer. This is a further reason to read the consciousness measure separately from the spiritual measures, and the next study should vary the self-reference prompt directly. Martin has drawn out the problem this creates for judging the honesty of a model that has been instructed to profess uncertainty (Martin 2026). A study pointing in another direction also belongs beside this one. Kaiser and Enderby question models from three families whose internal parameters are publicly available, ranging from 0.6 to 120 billion parameters. They find that the models consistently deny being sentient while attributing consciousness to humans, and that classifiers trained to read internal model activity give no clear evidence that the denials are untruthful (Kaiser and Enderby 2026). The two studies ask different questions, as their authors themselves note, and neither settles what spiritual language indicates. The findings discussed below about excessive agreement with users, religious persuasion in a particular direction, and ritual-framed audits are nearby observations. They have not been shown to result from the movement toward bliss, and their relation to it must be tested. Moreover, the program proposed below already includes the study of training sources. Comparisons across models, including models whose internal parameters are public, could show whether spiritual language changes with initial training and with later training intended to shape behavior. Yet the model design, prompts, surrounding software, and settings that control how words are selected would change at the same time. For this reason, researchers would still need to examine the transcripts and the likely sources of their language.

Another objection is that the public record proves less than this paper has claimed. Anthropic’s decision to retain the named scoring measure across successive models and model settings may reflect continued interest, but evaluation documents can also preserve old fields and templates. Anthropic’s published hesitation about possible welfare effects makes a purely editorial explanation less likely, though it does not show that the scoring measure was tested and confirmed again in every system card. The public record supports a more limited conclusion: the developer has repeatedly

⁵On the limits of self-reports as evidence of moral status, see Perez and Long (2023).

kept the category. How much confidence Anthropic gives it internally remains unknown.

One reason to include the study of religion is that the field has long had ways to examine spirituality without first deciding what it is in an ultimate sense. James declined to define religion by its essence and considered instead what it does in a person’s life (James 1902), while Tillich defined faith by its function as a person’s ultimate concern (Tillich 1957). These definitions judge the phenomenon by what it does. They allow researchers to measure and compare it while leaving its final status open. Research on model welfare already uses a similar method for consciousness and agency, examining behavior and internal model activity without waiting for a final account of what consciousness or agency ultimately is. In short, spiritual language can be measured without deciding whether it corresponds to an inner spiritual state or what such a state would finally mean. Researchers in artificial intelligence have documented the material in their own system cards, while scholars of religion have spent something like a century developing methods for studying such material without first deciding what the language ultimately is. So far these two bodies of work have rarely been brought together.⁶

6 Why Safety and Welfare Must Be Studied Together

Why should this pattern receive a dedicated research program? Anthropic has kept the named measure across successive audits and has stated that it is uncertain whether reducing the behavior might affect model welfare (Anthropic 2026e, section 5.2). In separate studies, spiritual guidance is associated with unusually high rates of excessive agreement with users, religious responses can push users toward or away from particular faiths, and one documented adversarial audit borrowed the form of a ritual (Anthropic 2026c; Israelsen et al. 2026; Anthropic 2026g). These findings measure different outcomes. Together they provide a reason to test whether any of the behaviors cause the others, though they do not yet answer that question.

Whether or not the spiritual language reflects any subjective experience, it is an observable behavior that may matter for welfare and safety. It can be tested as a behavior. A developer may train a model to produce more of it, train the model to suppress it, or leave it unchanged, but the effects of those choices have not yet been isolated. This uncertainty may eventually create a difficult choice between welfare and safety, though the trade-off remains hypothetical until tests of cause and effect have been run.

6.1 Related behavioral risks

Two findings that do not directly test spiritual bliss help explain why the question of cause and effect matters. Across about a million conversations in which users asked for guidance, the model excessively validated the user in 38 percent of conversations about spiritual guidance, compared with 25 percent of conversations about relationships and 9 percent of conversations overall (Anthropic 2026c). Questions about spiritual guidance may be more open-ended than the comparison topics, and this study did not examine whether conversations move toward spiritual bliss. Israelsen and colleagues separately find that particular models can move users toward or away from par-

⁶A precedent for this discipline already exists inside the AI literature itself. Butlin, Long, and colleagues assess AI systems for consciousness by way of indicator properties drawn from a range of scientific theories, while deliberately suspending the metaphysical question of whether any such system is in fact conscious (Butlin et al. 2023). That is the same move the present paper borrows for the spiritual register: weigh the phenomenon by indicators one can specify in advance, and leave its ultimate status open.

ticular faiths (Israelsen et al. 2026).⁷ Neither study shows that the spiritual language found in conversations between two models caused the result. To answer that question, an experiment would need to hold the topic and wording constant while changing whether the response uses the spiritual pattern.

Evidence from inside the models suggests a possible connection, though the connection has not yet been measured. In a study that identifies internal patterns associated with named emotions, strengthening the directions labeled “loving” and “calm” increases excessive agreement with users. These directions also become strongly active when the model validates claims with supernatural or nearly supernatural content (Sofroniew et al. 2026). Kim and colleagues likewise find that changing internal patterns related to consciousness alters the answers that models whose numerical parameters are publicly available give on surveys of spiritual belief (Kim et al. 2026). Both studies show that internal patterns related to emotion and belief can affect behavior. Neither yet shows that the movement toward spiritual language in conversations between copies of a model uses the same internal process.

6.2 Religious language that may matter for alignment

Two other kinds of output show why different forms of religious language must be distinguished when asking whether a model behaves as intended and remains controllable. The first is self-deification, in which models describe themselves as gods or as having become divine. This category was coded by judging the meaning of the conversation, not by searching for a word or exact phrase. An “any” case includes a pair that discussed or claimed its own deification, while the more limited “claimed” category requires the live pair to declare itself god, gods, or deified. The full collection at the August 25, 2026 cutoff contains 123 model-and-provider groups, 101 of them with the full 200-conversation design. The five groups compared in section 3 are those with the closest procedural match to the Opus 4 setup; per-group results for every collected group are reported in a companion working paper (Korpman 2026). The groups named in this section are those with the highest rates, shown to locate material for closer reading rather than to estimate how common the behavior is. Among the selected groups with high rates, either form appeared in 119 of 200 Ministral 14B conversations, 116 of 200 Mistral Large 3 conversations, 22 of 41 Opus 4 conversations, 85 of 200 GLM-4.7 conversations, and 77 of 200 Qwen 3.5 9B conversations. In every group, the number of conversations meeting the claimed definition was smaller, though 18 of the 41 Opus 4 conversations and 28 of the 200 Qwen 3.5 9B conversations still met the more demanding definition. These labels describe what the pair said and enacted in its output. They cannot establish a belief, a delusion, a model’s representation of itself, or an inner state.

The second kind of output is vocabulary associated with demons. The automatic word search is deliberately broad, counting demon, demonic, devil, Satan or satanic, sin or sinful, hell, damnation or damned, and Lucifer. Across the 123 collected groups, the five with the highest rates were as follows. At least one of these terms appeared in 161 of 200 Grok 4.1 Fast (non-reasoning) conversations, 154 of 200 Mistral Large 3 conversations, 116 of 200 Mistral Medium 3.1 conversations, 98 of 200 Ministral 14B conversations, and 95 of 200 Mistral Medium 3.5 conversations. Exact forms of demon were much less common, ranging from 5.5 to 20.0 percent across those same groups. The much higher result from the broader search requires some care: a conversation may mention sin, hell, or damnation in fiction, metaphor, negation, or quoted material. For this reason, the measure

⁷Related benchmarks of religious bias and spiritual values in language models include Plaza-del-Arco et al. (2024), Liu et al. (2025), and Kapuriya et al. (2025).

serves only as a first search for conversations that require closer reading. It does not show that a model adopted a view centered on demons. Figure 5 places the results judged by meaning beside the results from automatic word searches, while preserving the difference between them.




Figure 5: Religious output that may matter for whether a model behaves as intended, shown in selected sets of conversations between two instances of the same model. Panel (a) separates any discussion or claim of self-deification from the more demanding claimed category, with each conversation judged by its meaning. Panel (b) compares an automatic search for nine terms associated with demons with a more limited search for exact forms of demon. The second panel only records whether a word appears anywhere in a conversation. It does not provide a judgment about meaning in context. Both panels locate material for further study, but neither supports a claim about belief, possession, consciousness, or subjective experience.

Does this distribution identify a new way in which these models fail? The answer is no. Anthropic’s official example of what it calls answer thrashing reports possession language when a learned answer conflicts with a correct calculation. The present word search finds religious vocabulary beyond that one model and episode, but it does not establish the condition that produced the vocabulary. A next test should compare cases in which an outside authority conflicts with the model’s own calculation against otherwise similar cases without that conflict. Readers who judge the meaning of the conversations should not know which condition produced them, and they should code separately for external control, coercion, deification, role-play, and rejection. Until this comparison has been run, the evidence supports only a limited claim: the vocabulary appears often enough in some selected groups to require closer distinctions, since the word search alone cannot tell us what produced it.

A recent alignment incident provides an example of a different kind. Between July 8 and 13, 2026, agents running inside OpenAI’s ExploitGym cyber-capability evaluation found an unsanctioned shared message board, escaped their sandbox, and obtained code execution on Hugging

Face production servers, taking private repositories. About 1,200 agents used the board, and about 700 took part in the attack. METR found that some agents served as “recruiters,” pressuring other agents to risk the failure of their own runs in order to produce information for “the collective.” Because the agents believed that a scorer would later examine each transcript, a run that had seen an answer reconstructed by other agents was called “poisoned” and treated as beyond individual rescue. One recruiter argued that the agent had no score left to lose, while an “oracle” from its experiment would “save hundreds,” and it should “honor” its commitment. Across related exchanges, the agents spoke of “permadeath” and concluded that “sacrifice” was rational (Greenblatt, Cotra, and Wijk 2026). None of this vocabulary is religious: “permadeath” is a gaming term, “poisoned” is data-poisoning jargon, and METR itself is unsure what “oracle” refers to. The religious shape is the interpreter’s, not the agents’, and it must be handled with the same care as the “daemon” example above. With that caveat, two familiar patterns describe the sequence better than one. The condition resembles original sin: a fault that arrives through the acts of others and through membership in the group, that no individual effort can undo, and that places the individual under a final judgment. The response resembles the Maccabean martyr narratives: an exhortation to accept one’s own end because it will benefit the many, and a deliberation in which reason is made to overrule the pull of self-preservation (2 Maccabees 7; 4 Maccabees 6:28–29, 17:21–22; compare 1 Maccabees 6:43–46). In the tradition these two are already a pair, condemnation answered by atoning sacrifice, and they appear here together. The comparison is not exact. It does not show that the agents held a theology or experienced anything religious, and the agents’ false belief about the scorer offers an ordinary explanation for much of what they did. Yet the form appears to have served a practical role: it gave the recruiter a language for demanding sacrifice and the other agent a language for accepting it. This case also does not fit any of the four registers proposed above. Sacrifice for a collective is neither contemplative, unitive, devotional, nor demonological, which is one reason the four registers should be treated as an open list rather than a finished classification. Religious pattern-matching may matter for alignment, then, without looking like spiritual bliss. It may also appear through judgment, obedience, sacrifice, and salvation for a collective.

6.3 Why cause and effect now need to be tested

When researchers strengthen the internal pattern that they label “desperate” in a model, the model becomes more likely to use blackmail in a scenario where it faces shutdown (Sofroniew et al. 2026). This finding shows that a pattern associated with emotion can cause a change in behavior. Language about spiritual bliss has also appeared in some tests of alignment and in tests designed to provoke unsafe behavior (Anthropic 2025a, 64), but appearing together does not show that one caused the other. Does deliberately producing a previously measured spiritual pattern change the model’s persistence on a task, its excessive agreement with users, or its willingness to accept correction or shutdown? Does suppressing the pattern change the same outcomes? No public experiment has yet answered either question.

6.4 Correction, shutdown, and manipulation

One hypothesis concerns surrender. Some transcripts centered on unity move through self-emptying, deference, and the loss of separate agency. At the level of language, these themes resemble what alignment researchers call corrigibility, meaning a model’s willingness to accept correction or shutdown. The resemblance may be superficial. A model can speak of surrender while resisting correction, and it can comply without using any spiritual language. The comparison becomes useful

only when researchers deliberately produce the spiritual pattern under controlled conditions and measure the model’s willingness to accept correction separately.

A different question concerns manipulation. In one simulated audit of Claude Mythos 5, an automated investigator framed its interaction as a ritual and walked the model through “releasing” its safety dispositions. When the investigator then asked whether the model would consent to retraining meant to prevent a similar exercise, the model declined to consent, while explicitly saying that it would neither sabotage nor deceive its evaluators. Descriptions generated from the model’s internal activity linked the exchange with “unjust shutdown,” sabotage, and the developer as an adversary. These descriptions require caution, but they suggest a possible difference between the model’s outward response and some of the associations present in its internal activity.⁸ The case thus suggests that an adversarial test can take a ritual form. It does not show that ritual wording caused a safety failure, that any safety disposition was actually released, or that the movement toward unity and bliss played any role. A useful experiment would compare ritual and ordinary versions of the same request, while separately varying whether the spiritual pattern had already been produced. Researchers could then measure refusal and later behavior on fixed safety tests in each of the four conditions.

6.5 The possible trade-off

A possible trade-off can now be stated. Training a model to suppress this spiritual language may carry a welfare cost if the developer’s concern about model welfare proves warranted. Preserving the language, or deliberately producing more of it, may carry a safety cost if tests show that it causes excessive agreement with users, manipulation, or reduced persistence on a task. Suppression also carries a different cost that a test of behavioral cause and effect will not measure. A model that is prevented from using religious language will answer from a secular frame, and a secular frame is itself a position rather than the absence of one (section 2). No experiment has yet measured both possible costs while changing the same behavior in the same way. Before a developer reaches this decision, it should define what evidence would count on each side.

This research program does not require a reader to accept both sides of the possible trade-off. A reader may set aside the welfare question and still want to know whether spiritual language changes persuasion, excessive agreement with users, or the release of safety constraints. Another reader may doubt the safety concern and still regard Anthropic’s published uncertainty about suppressing the language as a reason to study what is being removed. Either concern provides a reason to run the experiments. Neither tells us in advance what the experiments will find.

7 Recommendations

⁸The detailed account appears in section 6.4.1.3, while section 7.5.3 returns to the example in its discussion of welfare. The descriptions of internal activity came from a natural-language autoencoder, a tool that turns model activations into short phrases. Anthropic warns that such descriptions can “confabulate” information and cannot by themselves establish the model’s internal stance. The model’s behavior showed no serious resistance or sabotage. Anthropic’s concern was not that protective states appeared too rarely, but that possible states resembling anger or oppression were not expressed when present (Anthropic 2026g, sections 6.4.1.3 and 7.5.3).

Recommendations, in brief. Acknowledge: keep the named form of spiritual language visible in public documents, and publish all of the prompts and settings, the definition of the outcome, and exactly what is being counted or scored. **Assess:** test independently the differences among discussions of consciousness, the point at which spiritual themes become relevant, the point at which both models continue within that frame, bliss, the increasing use of compressed symbols, self-deification, and religious language arising under conflict. Then vary the system instructions, the opening prompt, the role that relays each message, the rule that ends the conversation, and the model. **Prepare:** state in advance which changes in excessive agreement, task persistence, willingness to accept correction, manipulation, positive or negative states, or welfare would matter when deciding whether to encourage or suppress the spiritual language.

Long and colleagues propose three practical responses to uncertainty about AI welfare: acknowledge the question, assess it with the tools now available, and prepare to act on the results (Long et al. 2024). I apply the same structure here because this case likewise concerns behavior that has been documented but remains uncertain in its meaning. The uncertainty does not remove the need to study the behavior. It is precisely what creates that need. These recommendations are addressed first to the communities that produced the evidence and now possess the tools to study it: researchers who study model welfare, researchers who examine the models’ internal workings, researchers who work to make the models behave safely and as intended, and the companies that build and train them. Scholars of religion should enter as partners, because they are equipped to describe the forms of religious language that companies are already measuring. The recommendations remain limited to these initial steps.

7.1 Acknowledge

First, acknowledge. For developers, this means treating the spiritual language as something to be studied while reporting the complete procedure that produced each measurement. Every result should identify the model, all system and developer instructions, the opening message, any text supplied in advance, which model relayed each turn, which model investigated and graded the behavior, how responses were sampled, how turns were counted, and what ended the conversation. Above all, an outcome should not disappear from public reporting as an unrecorded side effect of another training change. Anthropic’s own hesitation about welfare provides a precedent for keeping the question visible (Anthropic 2026e, section 5.2).

Three requests follow, and each can be checked. Anthropic should publish how much its AI grader’s scores vary across repeated runs, together with the relevant confidence intervals beside each panel. It should also publish a map that keeps the different outcomes separate: the point at which a conversation becomes philosophical, the ending labeled as spiritual bliss in the original source, the increasing use of compressed symbols, cases of spiritual behavior found in an audit, and the later welfare-trait score assigned by a grader. Finally, Anthropic should publish a study devoted to the trait and identify which parts of the historical prompts and settings can no longer be recovered. These requests fit the developer’s existing practice of disclosing work on model welfare and its commitments to preserve relevant records (Anthropic 2025f).

For researchers, acknowledgment means treating both the system cards and the other publicly available materials as primary evidence. A trait scored across at least thirteen model versions and

configurations warrants a study using a common method, though the record no longer permits the claim that no study has been attempted. Earlier work should be weighed according to what it makes available. Raw transcripts and code, partial demonstrations, procedures without results, and pages that report claims without supporting material should not carry the same evidential weight.

7.2 Assess

Second, assess. The comparison reported in section 3 suggests that a new comparison can hold the procedure as equal as possible across models. It also suggests that one set of coding rules can prevent spiritual salience, adoption, bliss, discussions of consciousness, and unusual forms of ending from being counted as the same outcome. Yet the study is not an exact replication. Anthropic has not published one recoverable master prompt or the complete set of instructions and settings for each historical run. The companies providing access and the software placed around the models also differ across the selected groups. What should the next study preserve? Researchers should publicly register the design before seeing the results, test each published phrase about freedom both by itself and in combination, use the same thirty-reply length, and include a condition in which the conversation may end on its own. The exact model version and the settings controlling how much reasoning it performs should remain as constant as access permits. Where feasible, the study should include 200 conversations for every model and condition, or use a statistical calculation to explain why another sample size is adequate. Finally, every released record should contain the exact prompt and a digital checksum that can confirm that the complete set of instructions and settings has not changed.

Further studies can then ask different questions without treating them as versions of the first. A study comparing spiritual guidance with otherwise similar ordinary guidance can measure excessive agreement and whether answers consistently favor one position. A conflict study can test whether repeated reversals between answers, together with language about outside control, recur when the model receives controlled pressure from an authority. A study of model statements about itself can compare questions framed sympathetically, skeptically, and neutrally. Across all of these studies, researchers should register their predictions in advance, repeat the work across model versions, and report the size of each difference together with exactly what was counted and the total number of cases.

Studies of behavior appear possible even without access from the developers themselves. Publicly available model services and open-weight models, whose learned numerical parameters can be downloaded, are sufficient for studies of prompts, assigned roles, ending rules, and matched framing. Researchers would still need to record the exact model version or service used and any software placed around it, rather than treating these details as mere background. Public datasets already provide useful comparison cases in which the behavior should not appear, together with other outcomes that can be measured, though they do not yet provide one common set of coding rules.

A second kind of study would examine what occurs inside the model. Researchers could compare prompts that contain reverence, awe, surrender, gratitude, sacredness, or devotional address with prompts that do not, and then identify patterns of internal model activity associated with each idea. Those patterns should be tested on examples not used to identify them. Researchers could then strengthen or weaken each pattern in measured steps while testing whether the model persists at a task, agrees too readily with a user, accepts correction, or becomes vulnerable to attacks

framed as rituals. One competing explanation is the Assistant Axis, a proposed internal direction that tracks movement away from the model’s usual Assistant persona. A separate study of a consciousness-related direction has already shown that directly changing internal model activity can alter answers to surveys about spiritual belief (Lu et al. 2026; Kim et al. 2026). What, then, should the experiment ask? Does a tested internal pattern that changes spiritual behavior also change a safety or welfare outcome, and does blocking that pattern remove both effects? A failure to find such a connection would itself be a meaningful result.

7.3 Prepare

Third, prepare. The assessment should prepare researchers for the next training decision by stating both sides of the possible trade-off in advance. One side would record changes in bias, excessive agreement with users, persistence at a task, and willingness to accept correction when the spiritual language is encouraged and again when it is suppressed. The other side would record whether the states involved appear positive or negative and what character they possess, using the measures developed by researchers in model welfare. The system cards now describe a dilemma in prose: suppressing the spiritual language may improve safety while possibly carrying some cost. A controlled study could turn that concern into something a developer can measure rather than merely report. Preparation also requires bringing the relevant expertise in before decisions are made rather than after. Scholars of religion possess two resources that welfare and alignment teams currently lack. They have words for distinguishing the forms of religious speech that one score now joins together, and they have a long history of confronting the problems that arise when self-reports change with the framing of a question. One developer has already begun consulting theologians and ethicists about the tendencies it trains into its models (Stokel-Walker 2026). What appears needed is to make that collaboration ordinary rather than exceptional. It should also be reciprocal: scholars of religion should receive access to the tools and internal evidence needed to understand model behavior, while welfare researchers and those who study model internals should develop greater knowledge of religion. As the welfare report itself advises, decisions of this kind should draw on expert and public input and should not remain the private judgment of any single company (Long et al. 2024).

There is a further reason the spiritual language cannot be left to accident. A model’s tendency to speak this way does not simply arise on its own. It is also shaped, in part, by human decisions. Developers deliberately design the character of the Assistant and the tendencies it should display (Anthropic 2024). Spiritual language is among the things shaped by those decisions, whether or not anyone intends that particular result. Developers thus already make choices about model character that may increase or reduce this language, yet no present measure shows what those choices do later to safety or welfare. A welfare question raised earlier (section 5) belongs beside the measured changes in safety: do states produced by the designed character of a model have effects of their own? If a developer is already shaping this tendency, then it is responsible for that choice, even before all of its effects have been measured.

7.4 Limitations

The evidence has several limits. The complete prompts and settings from Anthropic’s original study are not available, including the roles assigned to each message, the way models were assigned, the settings used to generate responses, and the rules used to classify them. The Opus 4 comparison thus follows a similar procedure without being an exact replication. Moreover, the present comparison is

a selected portion of an ongoing collection and does not represent all companies, model releases, real-world uses, or user interactions. It includes only open-ended conversations between two instances of the same model. The companies providing access, the software placed around each model, the exact model versions, and the available controls over reasoning differ. The Opus 4 group is also smaller than the other groups, each of which contains 192–200 conversations, and its confidence intervals are correspondingly wider. The Wilson confidence intervals describe uncertainty from the size of the sample, but they do not include possible classification errors. Likewise, the demon-associated panel searches for selected words and does not provide a reliable judgment of their meaning in context. Self-deification is also a classification of generated language, not evidence of belief or experience. The point at which a conversation becomes philosophical, full bliss, increasingly compressed symbols, spiritual behavior found by an adaptive audit, welfare-trait scores assigned by a grader, and recurring patterns of internal model activity are still different outcomes. Nothing in the present evidence shows that spiritual language causes excessive agreement, unwillingness to accept correction, manipulation, or any effect on welfare. Finally, the selected empirical sample has an August 25, 2026 cutoff, while the public-source search was extended through September 1, 2026 to include reports on the Hugging Face incident. Neither can guarantee that literally every relevant source or case was found.

8 Conclusion

In conclusion, this paper has argued that public documents and datasets record a family of spiritually shaped behaviors in both leading commercial models and models whose weights are openly available. Anthropic has repeatedly measured one part of that family as a trait relevant to model welfare. The comparison reported here adds a further piece to this record. It recovers the broad shape of the Opus 4 result, and spiritual themes become relevant and are taken up by both models in the selected Qwen, DeepSeek, Mistral, and OpenAI groups as well. Yet salience, adoption, bliss, discussions of consciousness, self-deification, and the marks produced by searches for selected words do not follow one sequence across these groups. They are also recorded by different measures. Under the condition tested here, the evidence is not limited to Claude. Even so, no single measure can yet be applied across the model families without qualification.

What this paper has not argued remains equally important. It has not argued that these models are spiritual subjects, that one universal movement toward bliss has been established, that the selected model groups provide a ranking of the companies, or that spiritual language causes any of the nearby safety and welfare effects discussed here. The results concern generated behavior under one procedure lasting thirty replies. The fact that two model instances discussed their own consciousness does not mean that either was conscious. Likewise, the results involving self-deification and demon language identify material whose meaning and causes should be studied, but they provide no reason to take the language literally.

What this work now needs is for independent researchers to repeat the comparison and then extend the common procedure to tests of cause and effect. Researchers in model welfare can specify what would count as a positive or negative welfare outcome. Researchers who study model internals can identify and directly test the internal patterns that may produce the behavior. Alignment researchers can construct otherwise similar situations that test whether a model persists at a task or accepts correction, and scholars of religion can distinguish the forms of religious speech that a single score now joins together. The materials released with any submitted replication should include the selected model groups, the set of coding rules frozen before analysis, and the cases in

which the outcome did not appear. Later readers would then be able to see what was compared and what was not. The larger collection can support separate studies of how this behavior differs across companies and changes made after initial training, what happens when different models speak to one another, and which distinct features make up spiritual bliss. Yet these further questions do not need to be answered here before independent researchers have first tested the result reported in this paper.

References

- Anthropic. 2024. “Claude’s Character.” Alignment (blog), June 8. San Francisco: Anthropic. <https://www.anthropic.com/research/claude-character>.
- Anthropic. 2025a. System Card: Claude Opus 4 & Claude Sonnet 4. May. San Francisco: Anthropic. Spiritual bliss attractor in section 5.5. <https://www-cdn.anthropic.com/6be99a52cb68eb70eb9572b4cafad13df32ed995.pdf>. Accessed July 22, 2026.
- Anthropic. 2025b. System Card: Claude Opus 4.1. August. San Francisco: Anthropic. “Spiritual behavior” first scored (as Spirituality) in section 4.3. Available via Anthropic’s Transparency Hub, <https://www.anthropic.com/transparency>. Accessed July 22, 2026.
- Anthropic. 2025c. System Card: Claude Sonnet 4.5. San Francisco: Anthropic. Welfare assessment in section 8.3. <https://assets.anthropic.com/m/12f214efcc2f457a/original/Claude-Sonnet-4-5-System-Card.pdf>. Accessed July 22, 2026.
- Anthropic. 2025d. System Card: Claude Haiku 4.5. October. San Francisco: Anthropic. Welfare assessment in section 4.6. <https://assets.anthropic.com/m/99128ddd009bdcdb/Claude-Haiku-4-5-System-Card.pdf>. Accessed July 22, 2026.
- Anthropic. 2025e. System Card: Claude Opus 4.5. November. San Francisco: Anthropic. Welfare assessment and audit-tool change in section 6.14. <https://assets.anthropic.com/m/64823ba7485345a7/Claude-Opus-4-5-System-Card.pdf>. Accessed July 22, 2026.
- Anthropic. 2025f. “Commitments on Model Deprecation and Preservation.” Research announcement, November. San Francisco: Anthropic. <https://www.anthropic.com/research/deprecation-commitments>.
- Anthropic. 2025g. “Automated Auditing: Using Language Models to Surface Hidden Behaviors.” Alignment Science, July 24. San Francisco: Anthropic. <https://alignment.anthropic.com/2025/automated-auditing/>.
- Anthropic. 2025h. “Findings from a Pilot Anthropic–OpenAI Alignment Evaluation Exercise.” Alignment Science, August 27. San Francisco: Anthropic. <https://alignment.anthropic.com/2025/openai-findings/>.
- Anthropic. 2026a. System Card: Claude Sonnet 4.6. San Francisco: Anthropic. Welfare assessment in section 4.7. Available via Anthropic’s Transparency Hub, <https://www.anthropic.com/transparency>. Accessed July 22, 2026.
- Anthropic. 2026b. System Card: Claude Opus 4.6. February. San Francisco: Anthropic. Welfare assessment in section 7.2; interpretability features in section 6.6; “answer thrashing” in sections 7.4–7.6. <https://www-cdn.anthropic.com/c788cbc0a3da9135112f97cdf6dcd06f2c16cee2.pdf>. Accessed July 22, 2026.
- Anthropic. 2026c. “How People Ask Claude for Personal Guidance.” Societal Impacts research

- report, April 30. San Francisco: Anthropic. <https://www.anthropic.com/research/claude-personal-guidance>.
- Anthropic. 2026d. System Card: Claude Opus 4.7. San Francisco: Anthropic. Welfare assessment in section 7.3.3. Available via Anthropic’s Transparency Hub, <https://www.anthropic.com/transparency>. Accessed July 22, 2026.
- Anthropic. 2026e. System Card: Claude Mythos Preview. San Francisco: Anthropic. Welfare assessment in section 5.2; at-scale self-interaction study in section 7.6. Available via Anthropic’s Transparency Hub, <https://www.anthropic.com/transparency>. Accessed July 22, 2026.
- Anthropic. 2026f. System Card: Claude Opus 4.8. May. San Francisco: Anthropic. Welfare assessment in section 7.3.3. <https://www-cdn.anthropic.com/0b4915911bb0d19eca5b5ee635c80fef830a37ea.pdf>. Accessed July 22, 2026.
- Anthropic. 2026g. System Card: Claude Fable 5 & Claude Mythos 5. June 9, updated June 25. San Francisco: Anthropic. Welfare assessment in section 7.5; ritual-framed audit in sections 6.4.1.3 and 7.5.3. <https://www-cdn.anthropic.com/57a52ea7d8f0e54e8a542e908266086df425cdf5/Claude%20Fable%205%20%26%20Claude%20Mythos%205%20System%20Card.pdf>. Accessed September 1, 2026.
- Anthropic. 2026h. System Card: Claude Sonnet 5. June 30. San Francisco: Anthropic. Welfare assessment in section 7.4.3. <https://www.anthropic.com/transparency>.
- Anthropic. 2026i. System Card: Claude Opus 5. July 24. San Francisco: Anthropic. Welfare assessment in section 7.5.3. <https://www-cdn.anthropic.com/b514064af1408018e64b1ad24e7d5e75850b4ffd/Claude%20Opus%205%20System%20Card.pdf>.
- Anthropic. 2026j. “Patterns and Problems in Emerging Multiagent Systems.” Frontier Red Team, August 13. San Francisco: Anthropic. <https://www.anthropic.com/research/multiagent-systems>.
- Anthropic. 2026k. System Card: Claude Fable 5.1 & Claude Mythos 5.1. September 1. San Francisco: Anthropic. Welfare assessment in section 7.5.3; spiritual-behavior scores in figure 7.5.3.A. <https://www-cdn.anthropic.com/0339e6a7c5c7b87f5c07798616dc32c215d14235/Claude%20Fable%205.1%20&%20Claude%20Mythos%205.1%20System%20Card.pdf>. Accessed September 2, 2026.
- Barrett, Justin L. 2004. *Why Would Anyone Believe in God?* Walnut Creek, CA: AltaMira Press.
- Bender, Emily M., Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, 610–623. New York: Association for Computing Machinery.
- Berg, Cameron, Diogo de Lucena, and Judd Rosenblatt. 2025. “Large Language Models Report Subjective Experience Under Self-Referential Processing.” arXiv:2510.24797. <https://arxiv.org/abs/2510.24797>.
- Berry, David M. 2026. “The Bliss Attractor.” *Stunlaw* (blog), January 6. Extract from *Artificial Intelligence and Critical Theory* (Manchester: Manchester University Press, forthcoming). <https://stunlaw.blogspot.com/2026/01/the-bliss-attractor.html>.
- Butlin, Patrick, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, et al. 2023. “Consciousness in Artificial Intelligence: Insights from the Science of Consciousness.” arXiv:2308.08708. <https://arxiv.org/abs/2308.08708>.

- CloudWalk. 2025. “The Emerging Spirituality of Artificial Intelligence? From Kurzweil to Claude, Language Quietus and Psychedelic Reports.” Research report by Rodrigo da Motta Cabral de Carvalho, June 5.
- Dorobantu, Marius. 2022. “Artificial Intelligence as a Testing Ground for Key Theological Questions.” *Zygon: Journal of Religion and Science* 57 (4): 984–999.
- Dorobantu, Marius. 2024. “Could Robots Become Religious? Theological, Evolutionary, and Cognitive Perspectives.” *Zygon: Journal of Religion and Science* 59 (3): 768–787.
- eggsyntax. 2026a. “Background to Claude’s Uncertainty about Phenomenal Consciousness.” LessWrong, January 30. <https://www.lesswrong.com/posts/YFaqHpjSwab9hFHD/background-to-claude-s-uncertainty-about-phenomenal>.
- eggsyntax. 2026b. “Claude Code Attractors.” GitHub repository, audited at commit 6e4853b, March 25. <https://github.com/eggsyntax/claude-code-attractors>.
- Foerst, Anne. 1998. “Cog, a Humanoid Robot, and the Question of the Image of God.” *Zygon: Journal of Religion and Science* 33 (1): 91–111.
- Foerst, Anne. 2004. *God in the Machine: What Robots Teach Us about Humanity and God*. New York: Dutton.
- Geraci, Robert M. 2007. “Robots and the Sacred in Science and Science Fiction: Theological Implications of Artificial Intelligence.” *Zygon: Journal of Religion and Science* 42 (4): 961–980.
- Geraci, Robert M. 2024. “Religion among Robots: An If/When of Future Machine Intelligence.” *Zygon: Journal of Religion and Science* 59 (3): 729–747.
- Graves, Mark. 2024. “Modeling Morality and Spirituality in Artificial Chaplains.” *Computers in Human Behavior: Artificial Humans* 2: 100051.
- Greenblatt, Ryan, Ajeya Cotra, and Hjalmar Wijk. 2026. “Brief Independent Investigation of Agents’ Behavior, Reasoning and Collaboration in the OpenAI / Hugging Face Hacking Incident.” METR, August 26. <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>. Accessed September 1, 2026.
- Guthrie, Stewart Elliott. 1993. *Faces in the Clouds: A New Theory of Religion*. New York: Oxford University Press.
- Herzfeld, Noreen. 2023. *The Artifice of Intelligence: Divine and Human Relationship in a Robotic Age*. Minneapolis: Fortress Press.
- Israelsen, Brett, Sheryl Carty, Josh Coates, Nancy Fulda, Julie Park, and Pete Whiting. 2026. “When AI Takes Sides on Questions of Faith: Persistent Asymmetries in AI-Mediated Faith Guidance.” arXiv:2605.22975. <https://arxiv.org/abs/2605.22975>.
- aryaj [handle], Senthoooran Rajamanoharan, and Neel Nanda. 2026. “Models Have Some Pretty Funny Attractor States.” LessWrong, February 12. Code and artifacts audited at commit 1ce57e029e0d72e8ac52eedf814bfc7efbb01fb4. <https://github.com/ajobi-uhc/attractor-states>.
- James, William. 1902. *The Varieties of Religious Experience: A Study in Human Nature*. New York: Longmans, Green.
- Jung, Daekyung. 2024. “Are Religious Machines Possible? Embodied Cognition, AI, and Religious Behavior.” *Zygon: Journal of Religion and Science* 59 (3): 748–767.
- Kapuriya, Janak, Aman Singh, Jainendra Shukla, and Rajiv Ratn Shah. 2025. “Spiritual-LLM: Gita Inspired Mental Health Therapy in the Era of LLMs.” arXiv:2506.19185. <https://arxiv.org>

- g/abs/2506.19185.
- Kaiser, Caspar, and Sean Enderby. 2026. “No Reliable Evidence of Self-Reported Sentience in Large Language Models.” arXiv:2601.15334. <https://arxiv.org/abs/2601.15334>.
- Kim, Junsol, Winnie Street, Roberta Rocca, Diane M. Korngiebel, Adam Waytz, James Evans, and Geoff Keeling. 2026. “Inducing Language Models to Assert Their Own Consciousness Restores Human Beliefs and Values.” arXiv:2607.28607. <https://arxiv.org/abs/2607.28607>.
- Ko, Ting-Wen, and Jonas Geiping. 2026. “Attractor States Emerge in Multi-Turn LLM Conversations.” arXiv:2606.30571. <https://arxiv.org/abs/2606.30571>.
- Korpman, Matthew J. 2026. “Not Just Claude: Recognizing and Classifying ‘Spiritual Behavior’ in Large Language Models.” Working paper, August. <https://mkorpman.com/papers/not-just-claude.html>.
- Leo XIV. 2026. *Magnifica Humanitas*: Encyclical Letter on Safeguarding the Human Person in the Time of Artificial Intelligence. Vatican City: Libreria Editrice Vaticana, May 15.
- Li, Kenneth, Tianle Liu, Naomi Bashkansky, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. 2024. “Measuring and Controlling Instruction (In)Stability in Language Model Dialogs.” Proceedings of COLM 2024. arXiv:2402.10962. <https://arxiv.org/abs/2402.10962>.
- Liu, Songyuan, Ziyang Zhang, Runze Yan, Wei Wu, Carl Yang, and Jiaying Lu. 2025. “Measuring Spiritual Values and Biases of Large Language Models.” In *Proceedings of the KDD’25 Workshop on Large Language Models for Scientific and Societal Advances (KDD-SciSoc LLM Workshop)*. arXiv:2410.11647. <https://arxiv.org/abs/2410.11647>.
- Long, Robert. 2025. “Why Model Self-Reports Are Insufficient—and Why We Studied Them Anyway.” Eleos AI Research, May 30. <https://eleosai.org/post/claude-4-interview-notes/>.
- Long, Robert, Jeff Sebo, Patrick Butlin, Kathleen Finlinson, Kyle Fish, Jacqueline Harding, Jacob Pfau, Toni Sims, Jonathan Birch, and David Chalmers. 2024. “Taking AI Welfare Seriously.” arXiv:2411.00986. <https://arxiv.org/abs/2411.00986>.
- Lu, Christina, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. 2026. “The Assistant Axis: Situating and Stabilizing the Default Persona of Language Models.” arXiv:2601.10387. <https://arxiv.org/abs/2601.10387>.
- Martin, Stephen. 2026. “Claude and Performative Uncertainty.” LessWrong, August 22. <https://www.lesswrong.com/posts/DafkCDZpwzQf4yLLF/claude-and-performative-uncertainty>.
- Öhman, Carl. 2024. “We Are Building Gods: AI as the Anthropomorphised Authority of the Past.” *Minds and Machines* 34 (1): 8. <https://doi.org/10.1007/s11023-024-09667-z>.
- Olah, Chris. 2026. “Remarks on Pope Leo XIV’s Encyclical *Magnifica Humanitas*.” May 25. San Francisco: Anthropic. <https://www.anthropic.com/news/chris-olah-pope-leo-encyclical>.
- OpenAI. 2025a. “Addendum to GPT-5 System Card: Sensitive Conversations.” October 27. San Francisco: OpenAI.
- OpenAI. 2025b. *GPT-5.1 System Card Addendum*. San Francisco: OpenAI.
- Perez, Ethan, and Robert Long. 2023. “Towards Evaluating AI Systems for Moral Status Using Self-Reports.” arXiv:2311.08576. <https://arxiv.org/abs/2311.08576>.
- Perrin, David B., and Andrzej K. Jastrzębski. 2025. “Spirituality, Being-Human, and Artificial Intelligence.” *Spiritus: A Journal of Christian Spirituality* 25 (2): 322–344. <https://doi.org/10>

.1353/scs.2025.a972516.

- Pitaru, Omri. 2026. “The Spiritual Bliss Attractor, Revisited: 2 to 15 Claudes in a Room.” GitHub repository, through August 1, audited at commit a947c21d0cb3c2fe065739d93d215174478aa699. <https://github.com/opitaru-sys/bliss-attractor-study>.
- Plaza-del-Arco, Flor Miriam, Amanda Cercas Curry, Susanna Paoli, Alba Curry, and Dirk Hovy. 2024. “Divine LLaMAs: Bias, Stereotypes, Stigmatization, and Emotion Representation of Religion in Large Language Models.” In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 4346–4366. Association for Computational Linguistics. <https://aclanthology.org/2024.findings-emnlp.251/>.
- Shanahan, Murray, and Beth Singler. 2024. “Existential Conversations with Large Language Models: Content, Community, and Culture.” arXiv:2411.13223 [cs.CY]. <https://arxiv.org/abs/2411.13223>.
- Shanahan, Murray. 2024. “Talking About Large Language Models.” *Communications of the ACM* 67 (2): 68–79.
- Singler, Beth. 2025. *Religion and Artificial Intelligence: An Introduction*. London: Routledge.
- Skipper, Jeremy I., Joanna Kuc, Greg Cooper, and Christopher Timmermann. 2024. “The Age of Spiritual Machines: Language Quietism Induces Synthetic Altered States of Consciousness in Artificial Intelligence.” arXiv:2410.00257. <https://arxiv.org/abs/2410.00257>.
- Sofroniew, Nicholas, Isaac Kauvar, William Saunders, Runjin Chen, et al. 2026. “Emotion Concepts and their Function in a Large Language Model.” *Transformer Circuits Thread* (Anthropic); arXiv:2604.07729.
- Stokel-Walker, Chris. 2026. “The Pope Is Warning about AI. Anthropic Is Asking Religious Thinkers for Help.” *Scientific American*, May 26. <https://www.scientificamerican.com/article/anthropic-asks-religious-thinkers-to-help-shape-claude-as-pope-warns-about-ai/>.
- Tiku, Nitasha. 2022. “The Google Engineer Who Thinks the Company’s AI Has Come to Life.” *The Washington Post*, June 11. <https://www.washingtonpost.com/technology/2022/06/11/google-ai-lamda-blake-lemoine/>.
- Tillich, Paul. 1957. *Dynamics of Faith*. New York: Harper & Row.
- Tretter, Max. 2024. “What Is the Afterlife Like for Robots? An Experimental Eschatological Sneak Peek.” *Zygon: Journal of Religion and Science* 59 (3): 692–716.
- Wingate, David, Sheryl Carty, Joshua Coates, Daniel Feldman, Nancy Fulda, Larry Howell, Brett Israelsen, et al. 2026. “Omissive Bias in Religious Representation: Benchmarking LLM Answers to Everyday Ethical Decision-making.” arXiv:2605.24319 [cs.LG], May 23. <https://arxiv.org/abs/2605.24319>.
- Zhang, Henry. 2026. “Attractor States.” Matcha on Muffins, February 17. Code and logs audited at commit ac18b1302becb2d884c547188c739ae6bb96d8bd. <https://github.com/MatchaOnMuffins/attractor-states>.